

# Genome Center Clusters



# Introduction

“If you were plowing a field, which would you rather use: Two strong oxen or 1024 chickens?”

–Seymour Cray

# Introduction

- What is a cluster?

# Introduction

- What is a cluster?
  - A collection of machines that work together

# Introduction

- What is a cluster?
  - A collection of machines that work together
- Why use a cluster?

# Introduction

- What is a cluster?
  - A collection of machines that work together
- Why use a cluster?
  - Chickens are cheaper than oxen

# Introduction

- What is a cluster?
  - A collection of machines that work together
- Why use a cluster?
  - Chickens are cheaper than oxen
  - and easier to feed

# Introduction

- What is a cluster?
  - A collection of machines that work together
- Why use a cluster?
  - Chickens are cheaper than oxen
  - and easier to feed
  - ...but try making 1024 chickens move in the same direction



# Clusterable Jobs

- Parallel

# Clusterable Jobs

- Parallel
  - fine-grained parallel

# Clusterable Jobs

- Parallel
  - fine-grained parallel
  - course-grained parallel

# Clusterable Jobs

- Parallel
  - fine-grained parallel
  - course-grained parallel
  - embarrassingly parallel

# Clusterable Jobs

- Parallel
  - fine-grained parallel
  - course-grained parallel
  - embarrassingly parallel
- Serial

# Clusterable Jobs

- Parallel
  - fine-grained parallel
  - course-grained parallel
  - embarrassingly parallel
- Serial
  - Can still benefit from a cluster environment

# Available Resources

- genbeo

# Available Resources

- genbeo
- shiraz



# Available Resources

- genbeo
- shiraz
- apple

# Genbeo

- 19 nodes

# Genbeo

- 19 nodes
- Dual Opteron

# Genbeo

- 19 nodes
- Dual Opteron
- 4G memory per node

# Genbeo

- 19 nodes
- Dual Opteron
- 4G memory per node
- Gigabit ethernet interconnect

# Genbeo

- 19 nodes
- Dual Opteron
- 4G memory per node
- Gigabit ethernet interconnect
- 3.2TB storage

# Genbeo

- 19 nodes
- Dual Opteron
- 4G memory per node
- Gigabit ethernet interconnect
- 3.2TB storage
  - 800G User home dirs and mysql

# Genbeo

- 19 nodes
- Dual Opteron
- 4G memory per node
- Gigabit ethernet interconnect
- 3.2TB storage
  - 800G User home dirs and mysql
  - 800G Shared applications



# Genbeo

- 19 nodes
- Dual Opteron
- 4G memory per node
- Gigabit ethernet interconnect
- 3.2TB storage
  - 800G User home dirs and mysql
  - 800G Shared applications
  - 800G Shared filesystem

# Genbeo

- 19 nodes
- Dual Opteron
- 4G memory per node
- Gigabit ethernet interconnect
- 3.2TB storage
  - 800G User home dirs and mysql
  - 800G Shared applications
  - 800G Shared filesystem
  - 800G backups

# Shiraz

- 110 nodes

# Shiraz

- 110 nodes
- Dual processor, dual core Opteron (4 cores/node)

# Shiraz

- 110 nodes
- Dual processor, dual core Opteron (4 cores/node)
- 32 nodes with 8G, remainder with 4G

# Shiraz

- 110 nodes
- Dual processor, dual core Opteron (4 cores/node)
- 32 nodes with 8G, remainder with 4G
- Gigabit ethernet interconnect

# Shiraz

- 110 nodes
- Dual processor, dual core Opteron (4 cores/node)
- 32 nodes with 8G, remainder with 4G
- Gigabit ethernet interconnect
- 4.4TB attached storage

# Shiraz

- 110 nodes
- Dual processor, dual core Opteron (4 cores/node)
- 32 nodes with 8G, remainder with 4G
- Gigabit ethernet interconnect
- 4.4TB attached storage
  - 1.2T User storage



# Shiraz

- 110 nodes
- Dual processor, dual core Opteron (4 cores/node)
- 32 nodes with 8G, remainder with 4G
- Gigabit ethernet interconnect
- 4.4TB attached storage
  - 1.2T User storage
  - 1.6T shared databases (PDB, blastdb, ...)

# Shiraz

- 110 nodes
- Dual processor, dual core Opteron (4 cores/node)
- 32 nodes with 8G, remainder with 4G
- Gigabit ethernet interconnect
- 4.4TB attached storage
  - 1.2T User storage
  - 1.6T shared databases (PDB, blastdb, ...)
  - 1.6T Shared filesystem

# Apple

- 37 nodes

# Apple

- 37 nodes
- Dual PowerPC G4

# Apple

- 37 nodes
- Dual PowerPC G4
- 1-2G memory per node

# Apple

- 37 nodes
- Dual PowerPC G4
- 1-2G memory per node
- Gigabit ethernet interconnect

# Apple

- 37 nodes
- Dual PowerPC G4
- 1-2G memory per node
- Gigabit ethernet interconnect
- 1.6TB attached storage

# Apple

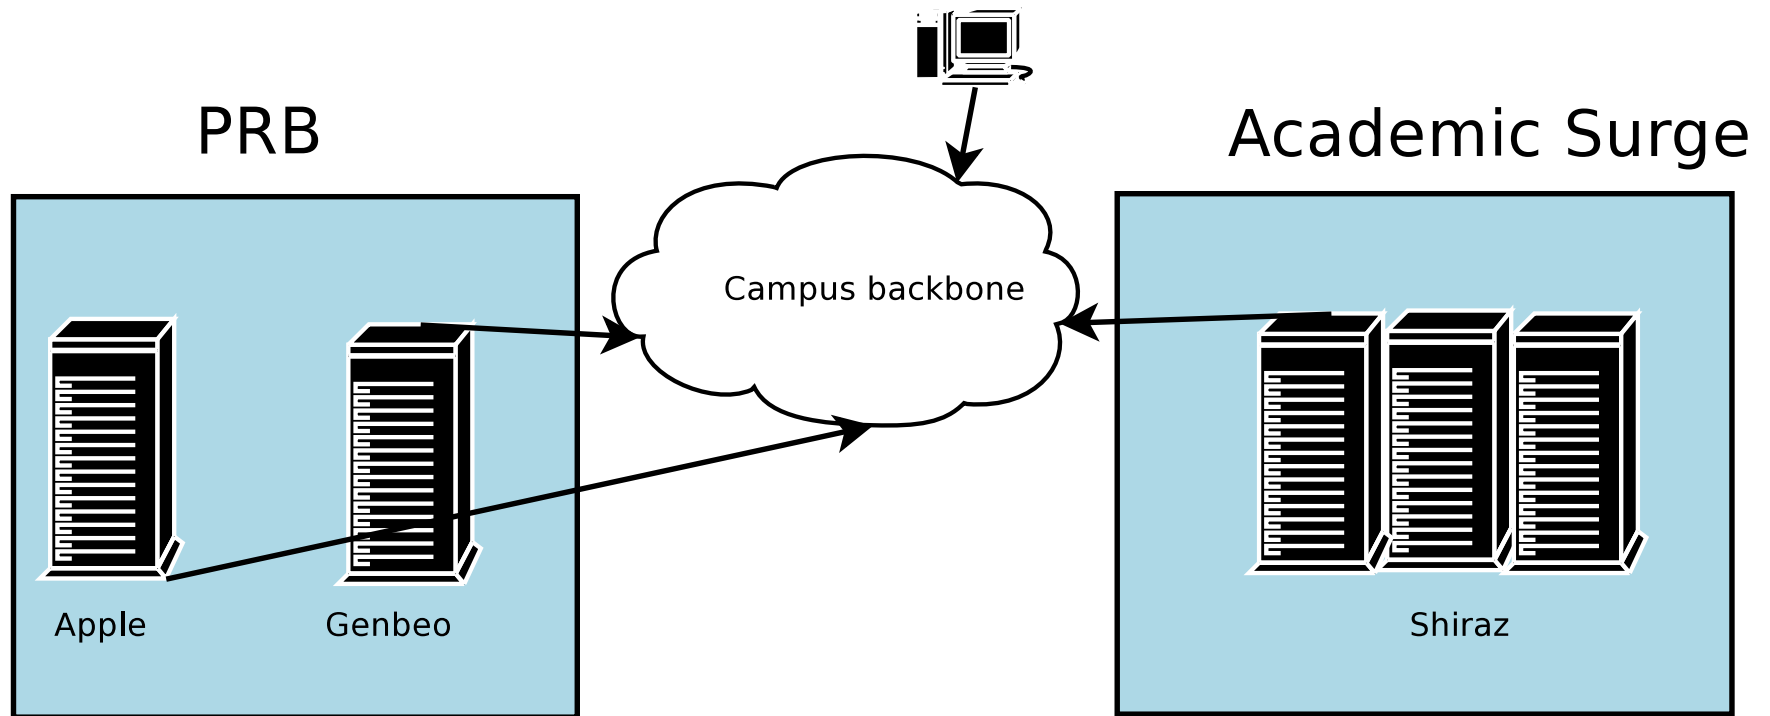
- 37 nodes
- Dual PowerPC G4
- 1-2G memory per node
- Gigabit ethernet interconnect
- 1.6TB attached storage
  - 800G User home directories



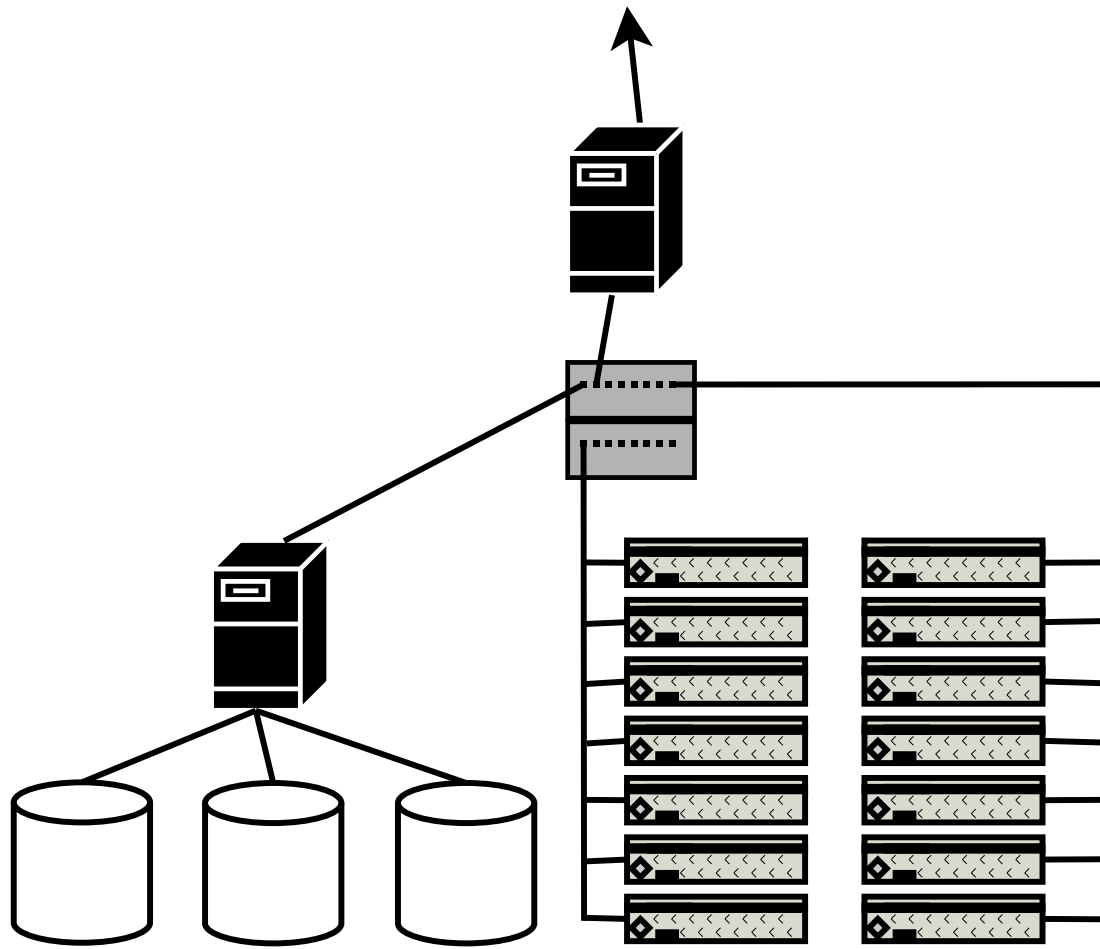
# Apple

- 37 nodes
- Dual PowerPC G4
- 1-2G memory per node
- Gigabit ethernet interconnect
- 1.6TB attached storage
  - 800G User home directories
  - 800G Mysql and postgres databases

# The Big Picture



# Anatomy of a Cluster



# Available Software

- Biological databases:

# Available Software

- Biological databases:
  - pdb

# Available Software

- Biological databases:
  - pdb
  - blastdb

# Available Software

- Biological databases:
  - pdb
  - blastdb
  - ..

# Available Software

## Compilers

- Gnu compiler suite



# Available Software

## Compilers

- Gnu compiler suite
- Pathscale floating license

# Available Software

## Compilers

- Gnu compiler suite
- Pathscale floating license
- Intel?

# Available Software

- biopython
- clustalw
- EMBOSS
- BLAST
- HMMER
- mrbayes
- t-coffee
- phylip
- ...

# Available Software

- R

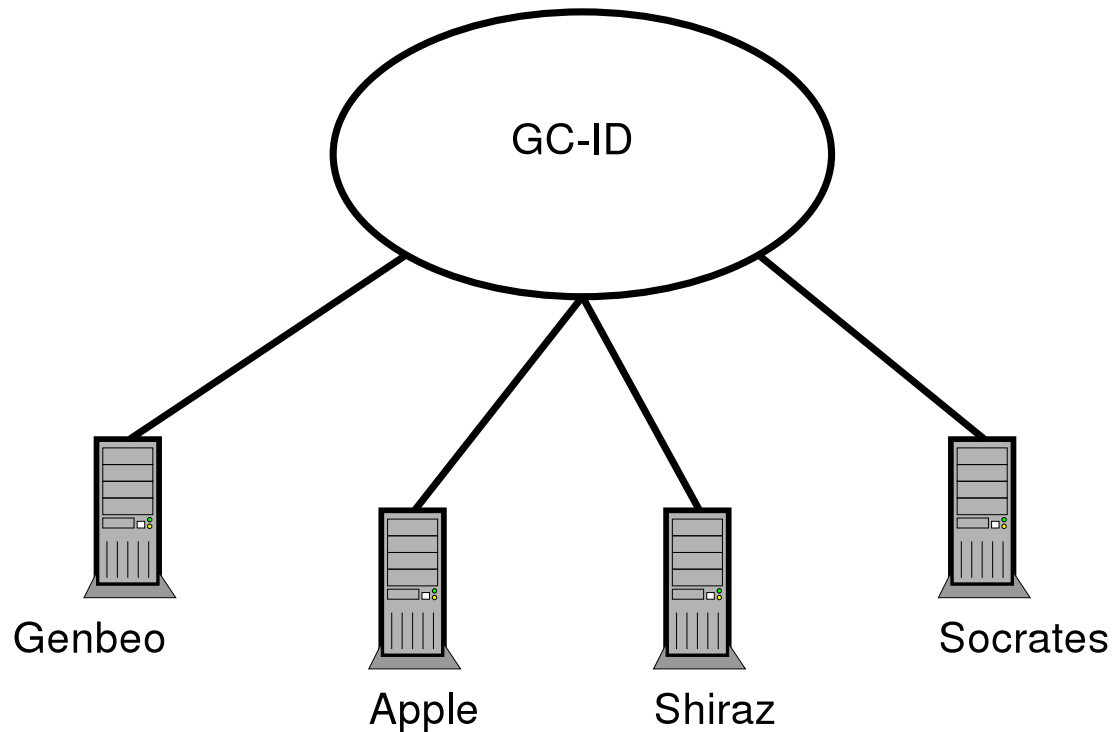
# Available Software

- R
- mysql

# Available Software

- R
- mysql
- Request your app!

# Your Cluster Account



Same username/password works on all hosts on which you have an account.

# Accessing the cluster

- Access using ssh



# Accessing the cluster

- Access using ssh
- Transfer files using scp/sftp/rsync

# Accessing the cluster

- Access using ssh
- Transfer files using scp/sftp/rsync
- Web interface

# Accessing the cluster: ssh

- `ssh username@genbeo.genomecenter.ucdavis.edu`

# Accessing the cluster: ssh

- `ssh username@genbeo.genomecenter.ucdavis.edu`
- Add `-X` option (capital!) to forward X connections

# Accessing the cluster: ssh

- `ssh username@genbeo.genomecenter.ucdavis.edu`
- Add `-X` option (capital!) to forward X connections
- Clients for windows:

# Accessing the cluster: ssh

- `ssh username@genbeo.genomecenter.ucdavis.edu`
- Add `-X` option (capital!) to forward X connections
- Clients for windows:
  - putty (free!)

# Accessing the cluster: ssh

- `ssh username@genbeo.genomecenter.ucdavis.edu`
- Add `-X` option (capital!) to forward X connections
- Clients for windows:
  - putty (free!)
  - SecureCRT

# Accessing the cluster: ssh

- `ssh username@genbeo.genomecenter.ucdavis.edu`
- Add `-X` option (capital!) to forward X connections
- Clients for windows:
  - putty (free!)
  - SecureCRT
  - ssh.com



# Accessing the cluster: ssh

- `ssh username@genbeo.genomecenter.ucdavis.edu`
- Add `-X` option (capital!) to forward X connections
- Clients for windows:
  - putty (free!)
  - SecureCRT
  - ssh.com
  - winscp

# Accessing the cluster: copying data

- scp

# Accessing the cluster: copying data

- scp
  - `scp file username@genbeo:`

# Accessing the cluster: copying data

- scp
  - `scp file username@genbeo:`
- rsync

# Accessing the cluster: copying data

- **scp**

- `scp file username@genbeo:`

- **rsync**

- `rsync -a directory username@genbeo:`

# Accessing the cluster: copying data

- scp

- `scp file username@genbeo:`

- rsync

- `rsync -a directory username@genbeo:`

- sftp

# Accessing the cluster: copying data

- scp

- `scp file username@genbeo:`

- rsync

- `rsync -a directory username@genbeo:`

- sftp

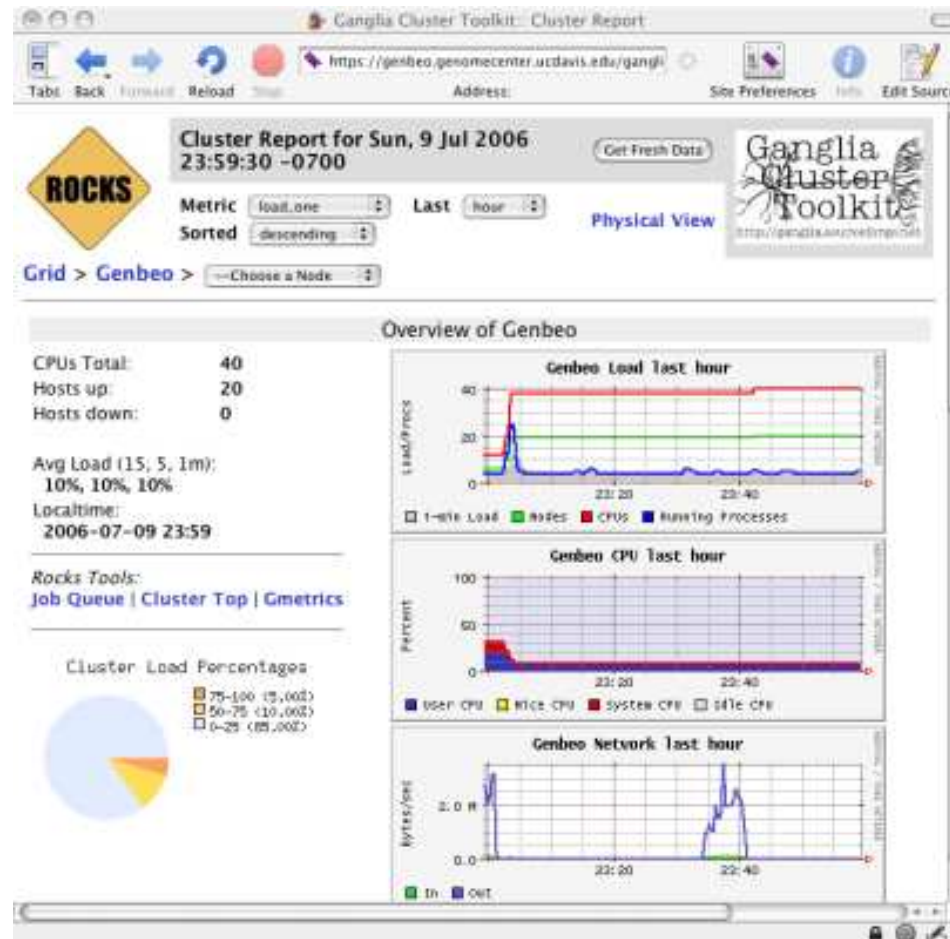
- ftp-like interface that works over ssh

# Cluster Web Interface





# Cluster Web Interface



# Cluster Web Interface

The screenshot shows a web browser window titled "Genbeo Job Queue" with the address bar displaying "https://genbeo.genomecenter.ucdavis.edu/gangli". The page header includes the title "Genbeo Job Queue" and a timestamp "Mon, 10 Jul 2006 00:00:30 -0700" with a link to "Physical Job Assignments". A yellow diamond-shaped warning sign with the word "ROCKS" is visible on the right. Below the header, there is a search bar labeled "Show only jobs for user:" with a "Go" button. The main content is a table of jobs with columns: Id, User, Processors, State, Name, Runtime, and TN. The table lists seven jobs, all in a "Running" state. At the bottom of the table, a summary line reads "7 Active Jobs, 19 of 38 Processors Active (50.00%)". A footer section contains a list of instructions for using the interface.

Genbeo Job Queue

Mon, 10 Jul 2006 00:00:30 -0700 Physical Job Assignments

Show only jobs for user:

Id	User	Processors	State	Name	Runtime	TN	Up   Down
869	smielke	1	Running	DNA_+_ions_ii_script	10 days, 13:31	6	
870	smielke	1	Running	DNA_+_ions_re_script	10 days, 13:2	6	
878	smielke	1	Running	DNA_+_ions_2M_script	10 days, 8:21	6	
903	cwu	4	Running	abfta.01	2 days, 14:56	6	
904	cwu	4	Running	abfta.02	2 days, 14:56	6	
905	cwu	4	Running	abftb.03	2 days, 14:56	6	
906	cwu	4	Running	abftb.04	2 days, 14:56	6	

7 Active Jobs, 19 of 38 Processors Active (50.00%)

- Click on a column header to sort by that field.
- Click on a job id for its nodes and history.
- TN is how old (in sect) the job information is.
- Note: Update button only shown if webserver is also the queue submit host.

# Sun Grid Engine

- What is a batch queue?

# Sun Grid Engine

- What is a batch queue?
  - Manage cluster resources

# Sun Grid Engine

- What is a batch queue?
  - Manage cluster resources
- Why use the batch queue?

# Sun Grid Engine

- What is a batch queue?
  - Manage cluster resources
- Why use the batch queue?
  - Share cluster resources

# Sun Grid Engine

- What is a batch queue?
  - Manage cluster resources
- Why use the batch queue?
  - Share cluster resources
  - You don't need to worry about when/where your jobs run

# Sun Grid Engine

- What is a batch queue?
  - Manage cluster resources
- Why use the batch queue?
  - Share cluster resources
  - You don't need to worry about when/where your jobs run
  - Submit a whole bunch of jobs and go home!



# SGE Terminology: slots

- A resource allocated to your job

# SGE Terminology: slots

- A resource allocated to your job
- We define one slot per CPU core

# SGE Terminology: slots

- A resource allocated to your job
- We define one slot per CPU core
- Can request number of slots when you submit job

# SGE Terminology: Parallel Environm

- Need to use a PE when you want more than one slot

# SGE Terminology: Parallel Environm

- Need to use a PE when you want more than one slot
- PE can specify start and stop programs (eg mpirun)

# SGE Terminology: Parallel Environm

- Need to use a PE when you want more than one slot
- PE can specify start and stop programs (eg mpirun)
- PEs available:

# SGE Terminology: Parallel Environm

- Need to use a PE when you want more than one slot
- PE can specify start and stop programs (eg mpirun)
- PEs available:
  - serial

# SGE Terminology: Parallel Environm

- Need to use a PE when you want more than one slot
- PE can specify start and stop programs (eg mpirun)
- PEs available:
  - serial
  - mpich



# SGE Terminology: Job Array

- Run the same job multiple times

# SGE Terminology: Job Array

- Run the same job multiple times
- submit/manage as a single job

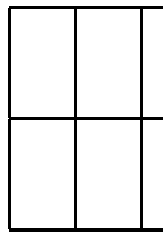
# SGE Terminology: Job Array

- Run the same job multiple times
- submit/manage as a single job
- Ideal for running the same program repeatedly with different input files or parameters

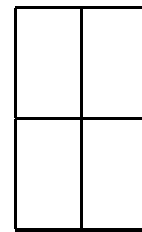
# Queue example

Job	Slots
<i>A</i>	6
<i>B<sub>1</sub></i>	1
<i>B<sub>2</sub></i>	1
<i>B<sub>3</sub></i>	1
<i>C</i>	4

Queue



Node1



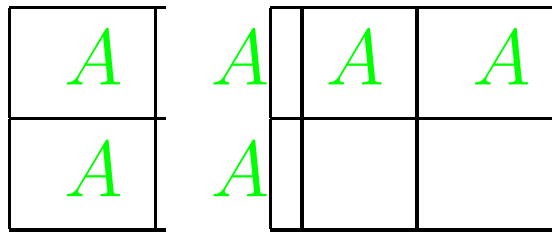
Node2

Three jobs waiting in queue...

# Queue example

Job	Slots
<i>A</i>	6
<i>B</i> <sub>1</sub>	1
<i>B</i> <sub>2</sub>	1
<i>B</i> <sub>3</sub>	1
<i>C</i>	4

Queue



Node1

Node2

Job *A* is scheduled

# Queue example

Job	Slots
<i>A</i>	6
<i>B<sub>1</sub></i>	1
<i>B<sub>2</sub></i>	1
<i>B<sub>3</sub></i>	1
<i>C</i>	4

Queue

<i>A</i>	<i>A</i>	<i>A</i>	<i>A</i>
<i>A</i>	<i>A</i>	<i>B<sub>1</sub></i>	

Node1

Node2

Job *B<sub>1</sub>* is scheduled

# Queue example

Job	Slots
<i>A</i>	6
<i>B<sub>1</sub></i>	1
<i>B<sub>2</sub></i>	1
<i>B<sub>3</sub></i>	1
<i>C</i>	4

Queue

<i>A</i>
<i>A</i>

Node1

<i>A</i>	<i>A</i>
<i>B<sub>1</sub></i>	<i>B<sub>2</sub></i>

Node2

Job *B<sub>2</sub>* is scheduled

# Queue example

Job	Slots
$B_1$	1
$B_2$	1
$B_3$	1
$C$	4

Queue


Node1

$B_1$	$B_2$

Node2

Job  $A$  finishes



# Queue example

Job	Slots
$B_1$	1
$B_2$	1
$B_3$	1
$C$	4

Queue

	$B_3$

Node1

$B_1$	$B_2$

Node2

Job  $B_3$  is scheduled

# Queue example

Job	Slots
$B_2$	1
$B_3$	1
$C$	4

Queue

	$B_3$		$B_2$

Node1

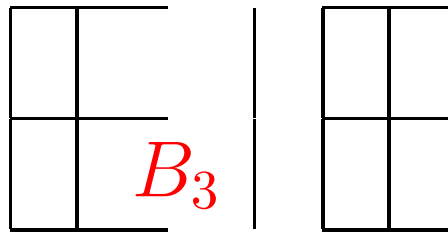
Node2

Job  $B_1$  finishes

# Queue example

Job	Slots
$B_3$	1
$C$	4

Queue



Node1

Node2

Job  $B_2$  finishes

# Queue example

Job	Slots
$B_3$	1
$C$	4

Queue

	$B_3$

Node1

$C$	$C$
$C$	$C$

Node2

Job  $C$  is scheduled

# SGE Commands

- qsub: submit jobs

# SGE Commands

- qsub: submit jobs
- qstat: get job status

# SGE Commands

- qsub: submit jobs
- qstat: get job status
- qdel: remove a job

# SGE Commands

- qsub: submit jobs
- qstat: get job status
- qdel: remove a job
- qlogin: interactive login



# SGE Commands: qsub

Use the `qsub` command to submit a batch job to the system

Simplest case:

```
$ qsub file.sh
```

```
Your job 929 ("file.sh") has been submitted.
```

# SGE Commands: qstat

Use the `-f` flag to see all jobs running...

```
$ qstat -f
```

queue	qtype	used/tot.	load_avg	arch	states
all.q@compute-0-1.local	BIP	4/4	1.83	lx26-amd64	
2455 0.60500 proAwt	cwu	r	06/21/2006 12:06:06		4
all.q@compute-0-10.local	BIP	0/4	0.00	lx26-amd64	d
all.q@compute-0-98.local	BIP	4/4	2.80	lx26-amd64	
2823 0.51386 ccr5_SCH_A	twang	r	07/07/2006 15:12:51		2
2865 0.50500 rungb5b	xjdeng	r	07/08/2006 17:37:06		1
2944 0.52905 g2l_ff03	zxwang	r	07/09/2006 22:07:21		1
all.q@compute-0-99.local	BIP	0/4	0.00	lx26-amd64	

# SGE Commands: qstat

...as well as those waiting to be run.

```
#####  
PENDING JOBS - PENDING JOBS - PENDING JOBS - PENDING JOBS - PENDING JOBS  
#####  
2947 0.52905 g2f_ff03    zxwang      qw      07/09/2006 22:11:04    20  
2948 0.52905 g2h_ff03    zxwang      qw      07/09/2006 22:11:55    20  
2949 0.52905 g2q_ff03    zxwang      qw      07/09/2006 22:12:42    20
```

# SGE Commands: qstat

Use the `-j <jobid>` flag to get more information about a job:

```
$ qstat -j 2947
```

```
job_number:                2947
```

```
exec_file:                 job_scripts/2947
```

```
submission_time:          Sun Jul  9 22:11:04 2006
```

```
...
```

# SGE Commands: qdel

- Use the `qdel` command to delete a previously scheduled job from the queue.

# SGE Commands: qdel

- Use the `qdel` command to delete a previously scheduled job from the queue.
- Note: if the job was running, you may still have to kill the processes by hand.

# SGE Commands: qdel

- Use the `qdel` command to delete a previously scheduled job from the queue.
- Note: if the job was running, you may still have to kill the processes by hand.
- The `-f` (force) option can sometimes be necessary to clean up jobs left behind, for example, if a node dies during the job.

# SGE Commands: qlogin

Use the `qlogin` command to schedule an interactive login.

- Default is one slot



# SGE Commands: qlogin

Use the `qlogin` command to schedule an interactive login.

- Default is one slot
- To allocate more slots, use the parallel environment serial and the number of slots  
`qlogin -pe serial 2`

# SGE Commands: qlogin

Use the `qlogin` command to schedule an interactive login.

- Default is one slot
- To allocate more slots, use the parallel environment serial and the number of slots  
`qlogin -pe serial 2`
- Two slots on genbeo will give you the whole node (4 on shiraz)

# SGE Commands: qlogin

Use the `qlogin` command to schedule an interactive login.

- Default is one slot
- To allocate more slots, use the parallel environment serial and the number of slots  
`qlogin -pe serial 2`
- Two slots on genbeo will give you the whole node (4 on shiraz)
- If enough slots are not available, `qlogin` will fail.

# SGE Example: BLAST

Single Threaded single app (simple) BLAST (Basic Local Alignment Search Tool)

SGE script: (sge\_blast.sh)

```
# Blast executable binary
BLAST=/opt/Bio/blast-2.2.14/bin/blastall

#$ -cwd
#$ -o blastout
#$ -e blasterr
#$ -q all.q
#$ -S /bin/bash

$BLAST $*
```

Lines with #\$ are command-line flags for SGE.

# SGE Example: BLAST

```
qsub -N serialBlast sge_blast.sh -p blastp -i query.aa  
-d drosoph.aa -e 1e-40 -o output.blast
```

- **-N Specify job name (serialBlast)**

# SGE Example: BLAST

```
qsub -N serialBlast sge_blast.sh -p blastp -i query.aa  
-d drosoph.aa -e 1e-40 -o output.blast
```

- -N Specify job name (serialBlast)
- -q Specify queue in which to run job (all.q)

# SGE Example: BLAST

```
qsub -N serialBlast sge_blast.sh -p blastp -i query.aa  
-d drosoph.aa -e 1e-40 -o output.blast
```

- -N Specify job name (serialBlast)
- -q Specify queue in which to run job (all.q)
- -S Specify interpreter to run script (/bin/bash)

# SGE Example: BLAST

```
qsub -N serialBlast sge_blast.sh -p blastp -i query.aa  
-d drosoph.aa -e 1e-40 -o output.blast
```

- -N Specify job name (serialBlast)
- -q Specify queue in which to run job (all.q)
- -S Specify interpreter to run script (/bin/bash)
- -o Specify file for stdout of job



# SGE Example: BLAST

```
qsub -N serialBlast sge_blast.sh -p blastp -i query.aa  
-d drosoph.aa -e 1e-40 -o output.blast
```

- -N Specify job name (serialBlast)
- -q Specify queue in which to run job (all.q)
- -S Specify interpreter to run script (/bin/bash)
- -o Specify file for stdout of job
- -cwd Execute job in the directory from which it was submitted

# SGE Example: mpi-BLAST

## Parallel version of BLAST using MPI SGE script: (sge\_mpiblast, sh)

```
# Basic SGE script to run mpiblast
```

```
MPIRUN=/opt/mpich/gnu/bin/mpirun
```

```
MPIBLAST=/opt/Bio/mpiblast-1.4.0/bin/mpiblast
```

```
#$ -cwd
```

```
#$ -o mpiblastout
```

```
#$ -e mpiblasterr
```

```
#$ -q all.q
```

```
#$ -S /bin/bash
```

```
$MPIRUN -np $NSLOTS -machinefile /$TMPDIR/machines $MPIBLAST $*
```

# SGE Example: mpi-BLAST

```
qsub -N mpiblast1 -pe mpich 4 sge_mpiblast.sh -p blastp  
-d swissprot -i query.aa -o mpiblast.out -e 1e-40
```

- **-pe** Specify parallel environment and number of slots

# SGE Example: R

- `-t 1-5:1`

# SGE Example: R

- `-t 1-5:1`
- `-hold_jid job-name-or-id`

# SGE Example: qlogin

# Things to do

- Use the scheduler!

# Things to do

- Use the scheduler!
- Checkpoint your job



# Things to do

- Use the scheduler!
- Checkpoint your job
- Make use of local storage on the nodes for intermediate results

# Things NOT to do

- Run jobs on the head node

# Things NOT to do

- Run jobs on the head node
- Many simultaneous writes to network filesystem

# Things NOT to do

- Run jobs on the head node
- Many simultaneous writes to network filesystem
- Go around scheduler and run directly on the nodes

# The end

- Get an account: <http://genomecenter.ucdavis.edu/>, click on "Administration"
- Documentation and example scripts:  
<http://wiki.genomecenter.ucdavis.edu/bioinformatics/>
- Questions?